

Post-Integration cleaning

	Missing_Count	Missing_Percentage
...		
EMOTIONAL_STATUS	3362596	52.098265
BODILY_INJURY	3362553	52.097598
COMPLAINT	3362546	52.097490
CROSS STREET NAME	2913209	45.135703
AGE_GROUP	1834119	28.416859
ON STREET NAME	1333227	20.656306
VEHICLE_ID	871637	13.504678
LONGITUDE	659594	10.219397
LATITUDE	659071	10.211294
LOCATION	637218	9.872715
PERSON_ID	630034	9.761410
POSITION_IN_VEHICLE	630015	9.761116
PERSON_SEX	630015	9.761116
CRASH_DATE	630015	9.761116
EJECTION	630015	9.761116
PED_ROLE	630015	9.761116
UNIQUE_ID	630015	9.761116
PERSON_TYPE	630015	9.761116
PERSON_INJURY	630015	9.761116
PERSON_AGE	630015	9.761116
CRASH_TIME	630015	9.761116
CRASH TIME	0	0.000000
BOROUGH	0	0.000000
ZIP_CODE	0	0.000000
CRASH DATE	0	0.000000

We dropped **EMOTIONAL_STATUS**, **BODILY_INJURY**, and **COMPLAINT** because each had over 50% missing values. Retaining them would weaken analysis and introduce bias, with no

reliable imputation strategy available. And we will drop the cross street name too as the percentage of missing values in it is high and this column will not provide informative insights.

```
df_integrated = df_integrated.drop(columns=['CRASH_DATE', 'CRASH_TIME'])
```

We dropped these 2 columns from the integrated table (table resulted from joining the 2 datasets) as they were redundant and there are other 2 columns that represent the same data.

```
CRASH DATE [class 'pandas._libs.tslibs.timestamps.Timestamp']
CRASH TIME [class 'datetime.time']
.. BOROUGH [class 'str']
ZIP CODE [class 'str' <class 'int'>]
LATITUDE [class 'float']
LONGITUDE [class 'float']
LOCATION [class 'float' <class 'str'>]
ON STREET NAME [class 'str' <class 'float'>]
NUMBER OF PERSONS INJURED [class 'float']
NUMBER OF PERSONS KILLED [class 'float']
NUMBER OF PEDESTRIANS INJURED [class 'int']
NUMBER OF PEDESTRIANS KILLED [class 'int']
NUMBER OF CYCLIST INJURED [class 'int']
NUMBER OF CYCLIST KILLED [class 'int']
NUMBER OF MOTORIST INJURED [class 'int']
NUMBER OF MOTORIST KILLED [class 'int']
CONTRIBUTING FACTOR VEHICLE 1 [class 'str']
CONTRIBUTING FACTOR VEHICLE 2 [class 'str']
COLLISION_ID [class 'int']
VEHICLE TYPE CODE 1 [class 'str']
VEHICLE TYPE CODE 2 [class 'str']
year [class 'int']
month [class 'pandas._libs.tslibs.period.Period']
UNIQUE_ID [class 'float']
PERSON_ID [class 'str' <class 'float'>]
PERSON_TYPE [class 'str' <class 'float'>]
PERSON_INJURY [class 'str' <class 'float'>]
VEHICLE_ID [class 'float']
PERSON_AGE [class 'float']
EJECTION [class 'str' <class 'float'>]
POSITION_IN_VEHICLE [class 'str' <class 'float'>]
PED_ROLE [class 'str' <class 'float'>]
PERSON_SEX [class 'str' <class 'float'>]
AGE_GROUP [class 'str' <class 'float'>]
```

Identified mixed data types across key columns (e.g., ZIP_CODE, LOCATION, PERSON_ID) requiring standardization before analysis.

```
[144] ✓ 2s
▶ # Convert ZIP_CODE to string to standardize mixed int/str values
df_integrated['ZIP CODE'] = df_integrated['ZIP CODE'].astype(str)

# Fill missing injury counts with 0 and enforce integer type
df_integrated['NUMBER OF PERSONS INJURED'] = df_integrated['NUMBER OF PERSONS INJURED'].fillna(0).astype(int)

# Convert UNIQUE_ID from float to nullable integer type for consistency
df_integrated['UNIQUE_ID'] = df_integrated['UNIQUE_ID'].astype('Int64')

# Impute PERSON_AGE with median and convert to integer for safe analysis
df_integrated['PERSON_AGE'] = df_integrated['PERSON_AGE'].fillna(df_integrated['PERSON_AGE'].median()).astype(in
```

The justification is in the comments.

	Missing_Count	Missing_Percentage
...		
EMOTIONAL_STATUS	3362596	52.098265
BODILY_INJURY	3362553	52.097598
COMPLAINT	3362546	52.097490
CROSS STREET NAME	2913209	45.135703
AGE_GROUP	1834119	28.416859
ON STREET NAME	1333227	20.656306
VEHICLE_ID	871637	13.504678
LONGITUDE	659594	10.219397
LATITUDE	659071	10.211294
LOCATION	637218	9.872715
PERSON_ID	630034	9.761410
POSITION_IN_VEHICLE	630015	9.761116
PERSON_SEX	630015	9.761116
CRASH_DATE	630015	9.761116
EJECTION	630015	9.761116
PED_ROLE	630015	9.761116
UNIQUE_ID	630015	9.761116
PERSON_TYPE	630015	9.761116
PERSON_INJURY	630015	9.761116
PERSON_AGE	630015	9.761116
CRASH_TIME	630015	9.761116
CRASH TIME	0	0.000000
BOROUGH	0	0.000000
ZIP CODE	0	0.000000
CRASH DATE	0	0.000000
month	0	0.000000
VEHICLE TYPE CODE 2	0	0.000000
year	0	0.000000
COLLISION_ID	0	0.000000
CONTRIBUTING FACTOR VEHICLE 2	0	0.000000
CONTRIBUTING FACTOR VEHICLE 1	0	0.000000

We observed that there are 6300015 crash records with no corresponding person records, so we filled these rows with "Unknown" for categorical and 0 for numerical.

```
[148] ✓ 12s
    # PERSON_AGE: numeric → impute missing with 0, cast to int
    df_integrated['PERSON_AGE'] = df_integrated['PERSON_AGE'].fillna(0).astype(int)

    # UNIQUE_ID: numeric identifier → impute missing with 0, cast to int
    df_integrated['UNIQUE_ID'] = df_integrated['UNIQUE_ID'].fillna(0).astype(int)

    # POSITION_IN_VEHICLE: categorical → impute missing with "Unknown"
    df_integrated['POSITION_IN_VEHICLE'] = df_integrated['POSITION_IN_VEHICLE'].astype(str).replace('nan', 'Unknown')

    # PERSON_SEX: categorical → impute missing with "Unknown"
    df_integrated['PERSON_SEX'] = df_integrated['PERSON_SEX'].astype(str).replace('nan', 'U')

    # EJECTION: categorical → impute missing with "Unknown"
    df_integrated['EJECTION'] = df_integrated['EJECTION'].astype(str).replace('nan', 'Unknown')

    # PED_ROLE: categorical → impute missing with "Unknown"
    df_integrated['PED_ROLE'] = df_integrated['PED_ROLE'].astype(str).replace('nan', 'Unknown')

    # PERSON_TYPE: categorical → impute missing with "Unknown"
    df_integrated['PERSON_TYPE'] = df_integrated['PERSON_TYPE'].astype(str).replace('nan', 'Unknown')

    # PERSON_INJURY: categorical → impute missing with "Unknown"
    df_integrated['PERSON_INJURY'] = df_integrated['PERSON_INJURY'].astype(str).replace('nan', 'Unknown')
```